Zeyuan Yang

https://faradawn.github.io/ | faradawn@uchicago.edu | (773)-231-3964

EDUCATION

The University of Chicago — Chicago, IL
Master of Science in Computer Science — Sept. 2023 - Jun. 2024
Bachelor of Science in Computer Science and Mathematics (GPA: 3.6/4.0) — Sept. 2020 - Jun. 2023

PROFESSIONAL EXPERIENCE

Ovative Group — Minneapolis, MN
Data Engineer — May. 2023 – Aug. 2023

- Enhanced data ingestion from **13.31s to 1.17s** by introducing Hive partition to a Hadoop storage.
- Reduced **4.7x** processing time by converting CSV to Parquet, creating the first **Delta Lake** for ETL data analytics.
- Managed a migration of 32 platforms from **GCP** to **Azure Databricks** to adopt **Apache Spark** ecosystem.

Seagate Technology — Chicago, IL
Software Engineer — Dec. 2021 – Jan. 2023

- Participated in writing a **C/C++** client for CORTX distributed cloud storage, enabling large-file upload.
- Undertook system testing and detected an inefficient 4KB-chunking algorithm, optimizing its speed by **2x**.
- Implemented a 16-node **Kubernetes** deployment by configuring RAM and CPU to enable scalability.

TikTok — Beijing, China
Software Engineer — Jun. 2021 – Sept. 2021

- Enhanced the runtime of a **JavaScript** Sudoku from **140s to 0.02s** for touch screens.
- Achieved a **9x** reduction in memory consumption by converting arrays to bitmasks.
- Created a full-stack **HTML-Django** scheduling website that saved visitors **10 minutes** of registration time.

Byto Inc. — Chicago, IL
Software Engineer — Jun. 2020 – Jun. 2021

- Developed a **Vue.js** job site adopted by **53** stores in Chicago. Agile software development cycle.
- Implemented a real-time chatroom with **NoSQL** and node.js backend that enhanced user engagement by **30%**.

RESEARCH & PROJECTS

Research: IONet – AI-powered cloud storage — University of Chicago
ML Researcher — Nov. 2021 – Present

- Built a deep learning cache prefetcher to achieve a hit rate of **87%**, surpassing the **40%** mark of Linux's default kernel.
- Designed a **TensorFlow** neural network to for storage center's I/O schedulers, reducing tail latency by **53%**.
- Wrote a multi-threaded I/O simulator in C that replayed Microsoft's traces, achieving **80-97%** accuracy.

IOS App: *Latin Garden* — Jan. 2021 – June. 2021

- Received a **4.9 rating** with **122** reviews on App Store**,** React Native, endorsed by Beijing Forestry University.

Enterprise Website: *US-Korea Conference 2022* — Jan. 2022 – Oct. 2022

- Developed an enterprise-level website streamlining paper submissions for **123** researchers.

SKILLS & INTERESTS

ACM-ICPC Programming Contest: 2023 Regional Silver Medal, 2022 Silver Medal.
Video creator: YouTube and Bilibili (5k subscribers).
Skills: C, C#, C++, Python, Java, JavaScript, R, SQL, PHP, Pytorch, GCP/AWS, data structure, database, Docker.
Courses taken: Operating Systems, Networks, Architecture, Machine Learning, Computer Vision, Statistics.